

# 15-418/618 Spring 2021

## Exercise 7

---

Assigned:	Fri., Apr. 2
Due:	Fri., Apr. 9, 11:00 pm

---

### Overview

This exercise is designed to help you better understand the lecture material and be prepared for the style of questions you will get on the exams. The questions are designed to have simple answers. Any explanation you provide can be brief—at most 3 sentences. You should work on this on your own, since that's how things will be when you take an exam.

You will submit an electronic version of this assignment to Canvas as a PDF file. For those of you familiar with the  $\text{\LaTeX}$  text formatter, you can download the template and configuration files at:

<http://www.cs.cmu.edu/~418/exercises/ex7.tex>  
<http://www.cs.cmu.edu/~418/exercises/config-ex7.tex>

Instructions for how to use this template are included as comments in the file. Otherwise, you can use this PDF document as your starting point. You can either: 1) electronically modify the PDF, or 2) print it out, write your answers by hand, and scan it. In any case, we expect your solution to follow the formatting of this document.

## Problem 1: Lock Implementations

For the following questions, we will compare different aspects of three lock implementations: test-and-test-and-set, ticket-based, and array-based lock.

- A. Releasing a lock requires a write. Order the three locks based on their interconnection traffic caused by the release. Briefly justify your ordering.

Test-and-test-and-set > ticket-based > array-based. Test and test and set requires a write and read per processor.  $O(P) * O(P) = O(P^2)$  interconnection traffic. Ticket-based requires one release broadcast per processor. When other processors observe broadcast they increment counter and check if it is their serve.  $O(P)$  interconnection traffic. Array-based doesn't need to broadcast release to all processors. Every processor check its own cache line.  $O(1)$  interconnection traffic.

- B. Compare the space requirements for the three lock implementations.

Test and test and set requires one int(lock variable) for all processor. Ticket-based requires one int(ticket variable) for all processor. Array-based requires one int per processor. Array-based > ticket-based = test and test and set.

- C. What is the advantage provided by the ticket-based and array-based locks versus a test-and-test-and-set (with or without exponential backoff)?

Ticket based gives advantage of reduced traffic. Only release needs to be broadcasted. Array based gives further reduced traffic. No need to broadcast at all but increases space requirements.

## Problem 2: Transactional Memory

The following code sequence attempts to use transactional memory to update a location in memory, and falls back on a lock when conflicts are detected; however, it has a flaw where updates can be lost. Please explain the flaw and how the code can be updated to prevent it. Recall that `*loc += val` is three instructions in x86 assembly.

```
void atomic_add(int* loc, int val)
{
    int result = xbegin();
    if (result == SUCCESS)
    {
        *loc += val;
        xend();
    }
    else
    {
        lock();
        *loc += val;
        unlock();
    }
}
```

I have thought two answers and not sure with both of them.

First of them we may need to insert stop the world flags both in transactional and fallback region to ensure abort function behaves correctly. But when I read specifications of xend function, it says that xend rolls back all values in case of terminated process, so I don't know if we have to use xabort function.

Second solution is to use while(true) loop around implementation to ensure that transaction happens for sure.